

In [20]:

```

1  import torch
2  import torchvision
3  import torch.nn as nn
4  import torchvision.transforms as transforms
5  import torchvision.datasets as dset
6  from torch.utils.data import DataLoader
7  from torch.autograd import Variable
8
9  '''
10 STEP 1: LOADING DATASET
11 '''
12
13 transform = transforms.Compose([transforms.ToTensor(), transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
14
15 trainset = torchvision.datasets.CIFAR10(root='./data', train = True, download = True, transform=transform)
16
17 testset = torchvision.datasets.CIFAR10(root='./data', train=False, download = True, transform=transform)
18
19 classes = ('plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
20
21 '''
22 STEP 2: MAKING DATASET ITERABLE
23 '''
24
25 batch_size = 100
26 n_iters = 3000
27 num_epochs = 20
28 #num_epochs = n_iters / (len(trainset) / batch_size)
29 #num_epochs = int(num_epochs)
30
31 print(len(trainset), num_epochs)
32
33 trainloader = torch.utils.data.DataLoader(trainset, batch_size = batch_size, shuffle = True, num_workers=4)
34
35 testloader = torch.utils.data.DataLoader(testset, batch_size = batch_size, shuffle = False, num_workers=4)
36
37 '''
38 STEP 3: CREATE MODEL CLASS
39 '''
40 class CNNModel(nn.Module):
41     def __init__(self):
42         super(CNNModel, self).__init__()
43
44         # Convolution 1
45         self.conv1 = nn.Conv2d(in_channels=3, out_channels=16, kernel_size=3, stride=1, padding=1)
46         self.relu1 = nn.ReLU()
47         self.batchNorm1 = nn.BatchNorm2d(16)
48
49         # Max pool 1
50         self.maxpool1 = nn.MaxPool2d(kernel_size=2)
51
52         # Convolution 2
53         self.conv2 = nn.Conv2d(in_channels=16, out_channels=32, kernel_size=3, stride=1, padding=1)
54         self.relu2 = nn.ReLU()
55         self.batchNorm2 = nn.BatchNorm2d(32)
56
57         # Max pool 2
58         self.maxpool2 = nn.MaxPool2d(kernel_size=2)
59

```

```
60     # Convolution 3
61
62     self.conv3 = nn.Conv2d(in_channels=32, out_channels=64, kernel_size=3, stride=1, padding=1)
63     self.relu3 = nn.ReLU()
64     self.batchNorm3 = nn.BatchNorm2d(64)
65
66     # Max pool 3
67     self.maxpool3 = nn.MaxPool2d(kernel_size=2)
68
69     # Convolution 4
70     self.conv4 = nn.Conv2d(in_channels=64, out_channels=128, kernel_size=3, stride=1, padding=1)
71     self.relu4 = nn.ReLU()
72     self.batchNorm4 = nn.BatchNorm2d(128)
73
74     # Max pool 4
75     self.maxpool4 = nn.MaxPool2d(kernel_size=2)
76
77     # Fully connected 1 (readout)
78     self.fc1 = nn.Linear(512, 512)
79
80     self.layer = nn.Sequential(
81         nn.Linear(512, 256),
82         nn.ReLU(),
83         nn.Linear(256, 128),
84         nn.ReLU(),
85         nn.Linear(128, 10)
86     )
87
88     def forward(self, x):
89         # Convolution 1
90         out = self.conv1(x)
91         out = self.relu1(out)
92         out = self.batchNorm1(out)
93
94         # Max pool 1
95         out = self.maxpool1(out)
96
97         # Convolution 2
98         out = self.conv2(out)
99         out = self.relu2(out)
100        out = self.batchNorm2(out)
101
102        # Max pool 2
103        out = self.maxpool2(out)
104
105        # Convolution 3
106        out = self.conv3(out)
107        out = self.relu3(out)
108        out = self.batchNorm3(out)
109
110        # Max pool 3
111        out = self.maxpool3(out)
112
113
114        # Convolution 4
115        out = self.conv4(out)
116        out = self.relu4(out)
117        out = self.batchNorm4(out)
118
119        # Max pool 4
120        out = self.maxpool4(out)
```

```

121
122     # Resize
123     # Original size: (100, 32, 7, 7)
124     # out.size(0): 100
125     # New out size: (100, 32*7*7)
126     out = out.view(out.size(0), -1)
127
128     # Linear function (readout)
129     out = self.fc1(out)
130     out = self.layer(out)
131
132     return out
133
134     '''
135 STEP 4: INSTANTIATE MODEL CLASS
136     '''
137
138     model = CNNModel().cuda()
139
140     #####
141     # USE GPU FOR MODEL #
142     #####
143
144     # device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
145
146
147
148     # model.to(device)
149
150     '''
151 STEP 5: INSTANTIATE LOSS CLASS
152     '''
153     criterion = nn.CrossEntropyLoss()
154
155
156     '''
157 STEP 6: INSTANTIATE OPTIMIZER CLASS
158     '''
159     learning_rate = 0.0015
160
161     optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate, weight_decay=0.002)
162
163

```

Files already downloaded and verified

Files already downloaded and verified

50000 20

In [21]:

```

1  '''
2  STEP 7: TRAIN THE MODEL
3  '''
4  iter = 0
5  for epoch in range(num_epochs):
6      for i, (images, labels) in enumerate(trainloader):
7
8          images = Variable(images).cuda()
9          labels = Variable(labels).cuda()
10
11         # Clear gradients w.r.t. parameters
12         optimizer.zero_grad()
13
14         # Forward pass to get output/logits
15         outputs = model(images)
16
17         # Calculate Loss: softmax --> cross entropy loss
18         loss = criterion(outputs, labels)
19
20         # Getting gradients w.r.t. parameters
21         loss.backward()
22
23         # Updating parameters
24         optimizer.step()
25
26         iter += 1
27
28         if iter % 200 == 0:
29             # Calculate Accuracy
30             correct = 0
31             total = 0
32             # Iterate through test dataset
33             for images, labels in testloader:
34                 #####
35                 # USE GPU FOR MODEL #
36                 #####
37                 #images = Variable(images).to(device)
38                 #labels = Variable(labels).to(device)
39                 images = Variable(images).cuda()
40                 labels = Variable(labels).cuda()
41
42                 # Forward pass only to get logits/output
43                 outputs = model(images)
44
45                 # Get predictions from the maximum value
46                 _, predicted = torch.max(outputs.data, 1)
47
48                 # Total number of labels
49                 total += labels.size(0)
50
51                 #####
52                 # USE GPU FOR MODEL #
53                 #####
54                 # Total correct predictions
55                 if torch.cuda.is_available():
56                     correct += (predicted.cpu() == labels.cpu()).sum()
57                 else:
58                     correct += (predicted == labels).sum()
59

```

```
60 accuracy = 100 * correct / total
61
62 # Print Loss
63 print('Iteration: {}. Loss: {}. Accuracy: {}'.format(iter, loss.item(), accuracy))
```

```
Iteration: 200. Loss: 1.0802453756332397. Accuracy: 54
Iteration: 400. Loss: 1.0591685771942139. Accuracy: 61
Iteration: 600. Loss: 1.2498900890350342. Accuracy: 63
Iteration: 800. Loss: 0.9499179124832153. Accuracy: 68
Iteration: 1000. Loss: 1.098762035369873. Accuracy: 70
Iteration: 1200. Loss: 0.7513238787651062. Accuracy: 69
Iteration: 1400. Loss: 0.84354567527771. Accuracy: 71
Iteration: 1600. Loss: 0.5185913443565369. Accuracy: 71
Iteration: 1800. Loss: 0.9087273478507996. Accuracy: 71
Iteration: 2000. Loss: 0.6466696262359619. Accuracy: 73
Iteration: 2200. Loss: 0.5459758639335632. Accuracy: 73
Iteration: 2400. Loss: 0.6114209890365601. Accuracy: 73
Iteration: 2600. Loss: 0.5416773557662964. Accuracy: 74
Iteration: 2800. Loss: 0.7845168113708496. Accuracy: 74
Iteration: 3000. Loss: 0.673625648021698. Accuracy: 75
Iteration: 3200. Loss: 0.48281776905059814. Accuracy: 74
Iteration: 3400. Loss: 0.6193227171897888. Accuracy: 74
Iteration: 3600. Loss: 0.5311362743377686. Accuracy: 75
Iteration: 3800. Loss: 0.7146061658859253. Accuracy: 73
Iteration: 4000. Loss: 0.6406211853027344. Accuracy: 76
Iteration: 4200. Loss: 0.4911739230155945. Accuracy: 76
Iteration: 4400. Loss: 0.6070870161056519. Accuracy: 76
Iteration: 4600. Loss: 0.5238867998123169. Accuracy: 76
Iteration: 4800. Loss: 0.3915778398513794. Accuracy: 76
Iteration: 5000. Loss: 0.5295023322105408. Accuracy: 77
Iteration: 5200. Loss: 0.43360623717308044. Accuracy: 76
Iteration: 5400. Loss: 0.40151968598365784. Accuracy: 77
Iteration: 5600. Loss: 0.3985865116119385. Accuracy: 76
Iteration: 5800. Loss: 0.6029594540596008. Accuracy: 76
Iteration: 6000. Loss: 0.42858999967575073. Accuracy: 77
Iteration: 6200. Loss: 0.32337379455566406. Accuracy: 78
Iteration: 6400. Loss: 0.3180921673774719. Accuracy: 78
Iteration: 6600. Loss: 0.32018253207206726. Accuracy: 78
Iteration: 6800. Loss: 0.5380628705024719. Accuracy: 78
Iteration: 7000. Loss: 0.5449371933937073. Accuracy: 78
Iteration: 7200. Loss: 0.4668041467666626. Accuracy: 78
Iteration: 7400. Loss: 0.3947407007217407. Accuracy: 79
Iteration: 7600. Loss: 0.33736830949783325. Accuracy: 78
Iteration: 7800. Loss: 0.3849654793739319. Accuracy: 78
Iteration: 8000. Loss: 0.47112786769866943. Accuracy: 78
Iteration: 8200. Loss: 0.38789236545562744. Accuracy: 79
Iteration: 8400. Loss: 0.4590197503566742. Accuracy: 79
Iteration: 8600. Loss: 0.2545468807220459. Accuracy: 79
Iteration: 8800. Loss: 0.40812888741493225. Accuracy: 79
Iteration: 9000. Loss: 0.35889342427253723. Accuracy: 80
Iteration: 9200. Loss: 0.362052321434021. Accuracy: 79
Iteration: 9400. Loss: 0.4707458019256592. Accuracy: 79
Iteration: 9600. Loss: 0.3114205300807953. Accuracy: 80
Iteration: 9800. Loss: 0.4438943564891815. Accuracy: 79
Iteration: 10000. Loss: 0.4702972173690796. Accuracy: 80
```

In [ ]:

1